

DataFrames

INTRODUCTION TO JULIA

James Fulton

Climate informatics researcher

Tabular data

	Day	Distance	Time	Raining
1	Wednesday	2000	14.99	true
2	Monday	5000	31.68	false
3	Thursday	3500	22.02	true
4	Tuesday	3000	17.25	true
5	Thursday	4500	25.47	false
6	Monday	5000	30.77	true

Tabular data

	Day	Distance	Time	Raining
1	Wednesday	2000	14.99	true
2	Monday	5000	31.68	false
3	Thursday	3500	22.02	true
4	Tuesday	3000	17.25	true
5	Thursday	4500	25.47	false
6	Monday	5000	30.77	true
	String	Int	Float	Bool

DataFrames

`using` DataFrames

```
# Create DataFrame
df = DataFrames.DataFrame(

)
```

DataFrames

`using` DataFrames

```
# Create DataFrame
```

```
df = DataFrame(  
    day = ["Wednesday", "Monday", "Thursday", "Tuesday", "Thursday", "Monday"]  
    distance = [2000, 5000, 3500, 3000, 4500, 5000]  
    time = [14.99, 31.68, 22.02, 17.25, 25.47, 30.77]  
    raining = [true, false, true, true, false, true]  
)
```

DataFrames

`using` DataFrames

```
# Create DataFrame
```

```
df = DataFrame(  
    day = ["Wednesday", "Monday", "Thursday", "Tuesday", "Thursday", "Monday"],  
    distance = [2000, 5000, 3500, 3000, 4500, 5000],  
    time = [14.99, 31.68, 22.02, 17.25, 25.47, 30.77],  
    raining = [true, false, true, true, false, true],  
)
```

DataFrames

```
println(df)
```

```
6×4 DataFrame
```

Row	day	distance	time	raining
	String	Int64	Float64	Bool
1	Wednesday	2000	14.99	true
2	Monday	5000	31.68	false
3	Thursday	3500	22.02	true
4	Tuesday	3000	17.25	true
5	Thursday	4500	25.47	false
6	Monday	5000	30.77	true

CSV files

- Comma separated variable
- Common format for tabular data

Inside `run.csv` :

```
day,distance,time,raining
Wednesday,2000,14.99,true
Monday,5000,31.68,false
Thursday,3500,22.02,true
Tuesday,3000,17.25,true
Thursday,4500,25.47,false
Monday,5000,30.77,true
```


Loading CSV files

```
using CSV
```

```
# Load the run data
```

```
file = CSV.File("run.csv")
```

```
# Convert the CSV file into the DataFrame
```

```
df = DataFrame(file)
```

- Cannot use `File("run.csv")` only `CSV.File("run.csv")`

Printing DataFrames

```
# Print the first 3 rows
println(first(df, 3))
```

```
3×4 DataFrame
 Row | day          distance  time    raining
     | String      Int64    Float64 Bool
-----|-----
  1 | Wednesday    2000    14.99    true
  2 | Monday       5000    31.68   false
  3 | Thursday     3500    22.02    true
```

Basic properties of DataFrames

```
# Print column names  
println(names(df))
```

```
["day", "distance", "time", "raining"]
```

```
# Print number of rows and columns  
println(size(df))
```

```
(6, 4)
```

Let's practice!

INTRODUCTION TO JULIA

Sorting and slicing data

INTRODUCTION TO JULIA

James Fulton

Climate informatics researcher

Selecting an element from the DataFrame

```
df_run = DataFrame(CSV.File("run.csv"))
println(df_run)
```

6×4 DataFrame

Row	day	distance	time	raining
	String	Int64	Float64	Bool
1	Wednesday	2000	14.99	true
2	Monday	5000	31.68	false
3	Thursday	3500	22.02	true
4	Tuesday	3000	17.25	true
5	Thursday	4500	25.47	false
6	Monday	5000	30.77	true

```
# df[rownum, colnum]
t = df_run[6, 3]
println(t)
```

30.77

```
# df[rowrange, colnum]
ts = df_run[5:6, 3]
println(ts)
```

[25.47, 30.77]

Selecting an element from the DataFrame

```
df_run = DataFrame(CSV.File("run.csv"))
println(df_run)
```

6×4 DataFrame

Row	day	distance	time	raining
	String	Int64	Float64	Bool
1	Wednesday	2000	14.99	true
2	Monday	5000	31.68	false
3	Thursday	3500	22.02	true
4	Tuesday	3000	17.25	true
5	Thursday	4500	25.47	false
6	Monday	5000	30.77	true

```
# df[rownum, colnum]
t = df_run[6, 3]
println(t)
```

30.77

```
# df[rowrange, colnum]
ts = df_run[end-1:end, 3]
println(ts)
```

[25.47, 30.77]

Selecting a column

```
df_run = DataFrame(CSV.File("run.csv"))
println(df_run)
```

6×4 DataFrame

Row	day	distance	time	raining
	String	Int64	Float64	Bool
1	Wednesday	2000	14.99	true
2	Monday	5000	31.68	false
3	Thursday	3500	22.02	true
4	Tuesday	3000	17.25	true
5	Thursday	4500	25.47	false
6	Monday	5000	30.77	true

```
# df[:, colnum]
distances = df_run[:, 2]
```

```
println(distances)
```

```
[2000, 5000, 3500, 3000, 4500, 5000]
```


Selecting a column

```
df_run = DataFrame(CSV.File("run.csv"))
println(df_run)
```

6×4 DataFrame

Row	day	distance	time	raining
	String	Int64	Float64	Bool
1	Wednesday	2000	14.99	true
2	Monday	5000	31.68	false
3	Thursday	3500	22.02	true
4	Tuesday	3000	17.25	true
5	Thursday	4500	25.47	false
6	Monday	5000	30.77	true

```
# df[:, colnum]
distances = df_run[:, 2]
```

```
# df[:, "colname"]
distances = df_run[:, "distance"]
```

```
# df.colname
distances = df_run.distance
```

```
println(distances)
```

```
[2000, 5000, 3500, 3000, 4500, 5000]
```

Selecting an element from the DataFrame

```
df_run = DataFrame(CSV.File("run.csv"))
println(df_run)
```

6×4 DataFrame

Row	day	distance	time	raining
	String	Int64	Float64	Bool
1	Wednesday	2000	14.99	true
2	Monday	5000	31.68	false
3	Thursday	3500	22.02	true
4	Tuesday	3000	17.25	true
5	Thursday	4500	25.47	false
6	Monday	5000	30.77	true

```
# df[rownum, colnum]
d = df_run[6, 2]
```

```
# df[rownum, "colname"]
d = df_run[6, "distance"]
```

```
# df.colname[rownum]
d = df_run.distance[6]
```

```
println(d)
```

```
5000
```

Slicing multiple columns

```
df_run = DataFrame(CSV.File("run.csv"))
println(df_run)
```

6×4 DataFrame

Row	day	distance	time	raining
	String	Int64	Float64	Bool
1	Wednesday	2000	14.99	true
2	Monday	5000	31.68	false
3	Thursday	3500	22.02	true
4	Tuesday	3000	17.25	true
5	Thursday	4500	25.47	false
6	Monday	5000	30.77	true

```
df_3cols = df_run[:, 1:3]
println(df_3cols)
```

6×3 DataFrame

Row	day	distance	time
	String	Int64	Float64
1	Wednesday	2000	14.99
2	Monday	5000	31.68
3	Thursday	3500	22.02
4	Tuesday	3000	17.25
5	Thursday	4500	25.47
6	Monday	5000	30.77

Selecting rows

```
df_run = DataFrame(CSV.File("run.csv"))
println(df_run)
```

6×4 DataFrame

Row	day	distance	time	raining
	String	Int64	Float64	Bool
1	Wednesday	2000	14.99	true
2	Monday	5000	31.68	false
3	Thursday	3500	22.02	true
4	Tuesday	3000	17.25	true
5	Thursday	4500	25.47	false
6	Monday	5000	30.77	true

```
# df[rownum, :]  
println(df_run[4, :])
```

DataFrameRow

Row	day	distance	time	raining
4	Tuesday	3000	17.25	true

Selecting multiple rows

```
df_run = DataFrame(CSV.File("run.csv"))
println(df_run)
```

6×4 DataFrame

Row	day	distance	time	raining
	String	Int64	Float64	Bool
1	Wednesday	2000	14.99	true
2	Monday	5000	31.68	false
3	Thursday	3500	22.02	true
4	Tuesday	3000	17.25	true
5	Thursday	4500	25.47	false
6	Monday	5000	30.77	true

```
# df[rowrange, :]  
println(df_run[2:4, :])
```

3×4 DataFrame

Row	day	distance	time	raining
	String	Int64	Float64	Bool
1	Monday	5000	31.68	false
2	Thursday	3500	22.02	true
3	Tuesday	3000	17.25	true

Sorting DataFrames

```
df_sort = sort(df_run, "time")
println(df_sort)
```

6×4 DataFrame

Row	day	distance	time	raining
	String	Int64	Float64	Bool
1	Wednesday	2000	14.99	true
2	Tuesday	3000	17.25	true
3	Thursday	3500	22.02	true
4	Thursday	4500	25.47	false
5	Monday	5000	30.77	true
6	Monday	5000	31.68	false

```
df_sort = sort(df_run, "time", rev=true)
println(df_sort)
```

6×4 DataFrame

Row	day	distance	time	raining
	String	Int64	Float64	Bool
1	Monday	5000	31.68	false
2	Monday	5000	30.77	true
3	Thursday	4500	25.47	false
4	Thursday	3500	22.02	true
5	Tuesday	3000	17.25	true
6	Wednesday	2000	14.99	true

Cheat sheet

Select a column

- `df[:, "colname"]`
- `df[:, colnum]`
- `df.colname`

Select a row

`df[rownum, :]`

Select multiple columns

`df[:, colnum1:colnum2]`

Select multiple rows

`df[rownum1:rownum2, :]`

Select a single value

- `df[rownum, "colname"]`
- `df[rownum, colnum]`
- `df.colname[rownum]`

Sorting by column

Ascending order

`sort(df, "colname")`

Descending order

`sort(df, "colname", rev=true)`

Let's practice!

INTRODUCTION TO JULIA

Descriptive statistics

INTRODUCTION TO JULIA

James Fulton

Climate informatics researcher

Describe function

```
# Summarize the runs DataFrame
println(describe(df_run))
```

4×7 DataFrame

Row	variable	mean	min	median	max	nmissing	eltype
	Symbol	Union{...}	Any	Union{...}	Any	Int64	DataType
1	day		Monday		Wednesday	0	String
2	distance	3833.33	2000	4000.0	5000	0	Int64
3	time	23.6967	14.99	23.745	31.68	0	Float64
4	raining	0.666667	false	1.0	true	0	Bool

Summary statistics on columns

`using` `Statistics`

Functions in `Statistics`:

- `mean()` - Calculate mean of array
- `median()` - Calculate median value of array
- `std()` - Calculate standard deviation of array values
- `var()` - Calculate variance of array values

```
# Calculate average of distance column  
average_distance = mean(df_run[:, "distance"])
```

Other builtin summary functions

- `sum()` - Calculate sum of array
- `minimum()` - Calculate minimum value in array
- `maximum()` - Calculate maximum value in array

```
total_distance = sum(df_run[:, "distance"]) # Returns 23000
```

```
minimum_distance = minimum(df_run[:, "distance"]) # Returns 2000
```

```
maximum_distance = maximum(df_run[:, "distance"]) # Returns 5000
```

Column operations

For columns `a` and `b` of DataFrame `df`

Operation	Scalar example	Array example
Addition	<code>df.a .+ 1</code>	<code>df.a .+ df.b</code> or <code>df.a + df.b</code>
Subtraction	<code>df.a .- 1</code>	<code>df.a .- df.b</code> or <code>df.a - df.b</code>
Multiplication	<code>2 .* df.a</code> or <code>2 * df.a</code>	<code>df.a .* df.b</code>
Division	<code>df.a ./ 2</code> or <code>df.a / 2</code>	<code>df.a ./ df.b</code>

Calculating run speed

```
# Convert distances to kilometers
distance_km = df_run.distance ./ 1000
```

```
# Convert run times to hours
time_hr = df_run.time ./ 60
```

```
println(distance_km)
println(time_hr)
```

```
[2.0, 5.0, 3.5, 3.0, 4.5, 5.0]
[0.25, 0.53, 0.37, 0.29, 0.42, 0.51]
```

```
6×4 DataFrame
 Row | distance      time      ...
     |      Int64    Float64    ...
-----|-----
  1 |      2000      14.99     ...
  2 |      5000      31.68     ...
  3 |      3500      22.02     ...
  4 |      3000      17.25     ...
  5 |      4500      25.47     ...
  6 |      5000      30.77     ...
```

Calculating run speed

```
# Convert distances to kilometers
distance_km = df_run.distance ./ 1000
```

```
# Convert run times to hours
time_hr = df_run.time ./ 60
```

```
# Run speed in km/hr
speeds = distance_km ./ time_hr

println(speeds)
```

```
[8.01, 9.47, 9.54, 10.43, 10.60, 9.75]
```

```
6×4 DataFrame
 Row | distance      time      ...
     |      Int64    Float64    ...
-----|-----
  1 |      2000      14.99      ...
  2 |      5000      31.68      ...
  3 |      3500      22.02      ...
  4 |      3000      17.25      ...
  5 |      4500      25.47      ...
  6 |      5000      30.77      ...
```

Column assignment

```
# Assign run speeds to new column named "speed"  
df_run[:, "speed"] = distance_km ./ time_hr
```

```
# Assign using dot form  
df_run.speed = distance_km ./ time_hr
```


Column assignment

```
println(df_run)
```

6×4 DataFrame

Row	day	distance	time	raining	speed
	String	Int64	Float64	Bool	Float64
1	Wednesday	2000	14.99	true	8.01
2	Monday	5000	31.68	false	9.47
3	Thursday	3500	22.02	true	9.54
4	Tuesday	3000	17.25	true	10.43
5	Thursday	4500	25.47	false	10.60
6	Monday	5000	30.77	true	9.75

Let's practice!
INTRODUCTION TO JULIA

Filtering

INTRODUCTION TO JULIA

James Fulton

Climate informatics researcher

Example

6×4 DataFrame

Row	day	distance	time	raining
	String	Int64	Float64	Bool
1	Wednesday	2000	14.99	true
2	Monday	5000	31.68	false
3	Thursday	3500	22.02	true
4	Tuesday	3000	17.25	true
5	Thursday	4500	25.47	false
6	Monday	5000	30.77	true

The filter function

6×4 DataFrame

Row	day	distance	time	raining
	String	Int64	Float64	Bool
1	Wednesday	2000	14.99	true
2	Monday	5000	31.68	false
...				

```
# Filter to Monday runs
```

```
df_monday = filter(row -> row.day=="Monday", df_run)
```

The filter function

```
println(df_monday)
```

```
6×4 DataFrame
```

Row	day	distance	time	raining
	String	Int64	Float64	Bool
1	Monday	5000	31.68	false
2	Monday	5000	30.77	true

Filtering on numerical columns

```
# Filter to shorter runs
df_short = filter(row -> row.distance <= 3000, df_run)
println(df_short)
```

Row	day	distance	time	raining
1	Wednesday	2000	14.99	true
2	Tuesday	3000	17.25	true

Filtering on boolean columns

```
# Filter to raining days
df_raining = filter(row -> row.raining, df_run)
println(df_raining)
```

Row	day	distance	time	raining
1	Wednesday	2000	14.99	true
2	Thursday	3500	22.02	true
3	Tuesday	3000	17.25	true
4	Monday	5000	30.77	true

Filtering on all comparisons

- `row.col == b` filter to where `row.col` equals `b`
- `row.col != b` filter to where `row.col` does not equal `b`
- `row.col > b` filter to where `row.col` is greater than `b`
- `row.col >= b` filter to where `row.col` is greater than or equal to `b`
- `row.col < b` filter to where `row.col` is less than `b`
- `row.col <= b` filter to where `row.col` is less than or equal to `b`
- `row.col` filter to where `row.col` is `true`

Further analysis

```
# Distance run in rain  
println(sum(df_raining.time))
```

```
13500
```

Let's practice!

INTRODUCTION TO JULIA

Final thoughts

INTRODUCTION TO JULIA

James Fulton

Climate informatics researcher

The exciting world of Julia

- Julia created 2012
- Expanding rapidly
- Stable version 1.0 released in 2018



Some things we've covered

- First steps

- `println(1+1)`

- `x = 0`

- `y = "fin"`

- Arrays

- `x = [1, 2, 3, 4, 5]`

- `x .* 2`

- `x[2:end]`

- Conditional statements

```
if enjoyment == 5
    println("Leave a rating?")
else
    println("Goodbye :)")
end
```

- Functions

```
function f(x)
    return x^2 + 2*x + 1
end
```

Some things we've covered

Broadcasting

```
x = [1, 2, 3]
y = f.(x)
```

Multiple dispatch

```
function f(x::Int64)
    return x^2 + 2*x + 1
end

function f(x::Bool)
    return x
end
```

Using Packages

```
import Statistics
using DataFrames
```

DataFrames

```
df = DataFrames(CSV.Files("run.csv"))
```

Congratulations!

INTRODUCTION TO JULIA